

```
# Set the CRAN mirror to RStudio's cloud server
options(repos = c(CRAN = "https://cloud.r-project.org"))

# Install the package
install.packages("corrplot")
```

The downloaded binary packages are in

/var/folders/kj/zb0ctvqs5zjbr\_x5cj8crxjh0000gn/T//RtmpmoMCiW/downloaded\_packages

```
library(leaps)
set.seed(1)
red_wine <- read.csv("winequality-red.csv", sep = ";")
white_wine <- read.csv("winequality-white.csv", sep = ";")
summary(red_wine)
```

fixed.acidity	volatile.acidity	citric.acid	residual.sugar
Min. : 4.60	Min. :0.1200	Min. :0.000	Min. : 0.900
1st Qu.: 7.10	1st Qu.:0.3900	1st Qu.:0.090	1st Qu.: 1.900
Median : 7.90	Median :0.5200	Median :0.260	Median : 2.200
Mean : 8.32	Mean :0.5278	Mean :0.271	Mean : 2.539
3rd Qu.: 9.20	3rd Qu.:0.6400	3rd Qu.:0.420	3rd Qu.: 2.600
Max. :15.90	Max. :1.5800	Max. :1.000	Max. :15.500
chlorides	free.sulfur.dioxide	total.sulfur.dioxide	density
Min. :0.01200	Min. : 1.00	Min. : 6.00	Min. :0.9901
1st Qu.:0.07000	1st Qu.: 7.00	1st Qu.: 22.00	1st Qu.:0.9956
Median :0.07900	Median :14.00	Median : 38.00	Median :0.9968
Mean :0.08747	Mean :15.87	Mean : 46.47	Mean :0.9967
3rd Qu.:0.09000	3rd Qu.:21.00	3rd Qu.: 62.00	3rd Qu.:0.9978
Max. :0.61100	Max. :72.00	Max. :289.00	Max. :1.0037
pH	sulphates	alcohol	quality
Min. :2.740	Min. :0.3300	Min. : 8.40	Min. :3.000
1st Qu.:3.210	1st Qu.:0.5500	1st Qu.: 9.50	1st Qu.:5.000
Median :3.310	Median :0.6200	Median :10.20	Median :6.000
Mean :3.311	Mean :0.6581	Mean :10.42	Mean :5.636
3rd Qu.:3.400	3rd Qu.:0.7300	3rd Qu.:11.10	3rd Qu.:6.000
Max. :4.010	Max. :2.0000	Max. :14.90	Max. :8.000

```
summary(white_wine)
```

fixed.acidity	volatile.acidity	citric.acid	residual.sugar
Min. : 3.800	Min. :0.0800	Min. :0.0000	Min. : 0.600
1st Qu.: 6.300	1st Qu.:0.2100	1st Qu.:0.2700	1st Qu.: 1.700
Median : 6.800	Median :0.2600	Median :0.3200	Median : 5.200
Mean : 6.855	Mean :0.2782	Mean :0.3342	Mean : 6.391
3rd Qu.: 7.300	3rd Qu.:0.3200	3rd Qu.:0.3900	3rd Qu.: 9.900
Max. :14.200	Max. :1.1000	Max. :1.6600	Max. :65.800

chlorides	free.sulfur.dioxide	total.sulfur.dioxide	density
Min. :0.00900	Min. : 2.00	Min. : 9.0	Min. :0.9871
1st Qu.:0.03600	1st Qu.: 23.00	1st Qu.:108.0	1st Qu.:0.9917
Median :0.04300	Median : 34.00	Median :134.0	Median :0.9937
Mean :0.04577	Mean : 35.31	Mean :138.4	Mean :0.9940
3rd Qu.:0.05000	3rd Qu.: 46.00	3rd Qu.:167.0	3rd Qu.:0.9961
Max. :0.34600	Max. :289.00	Max. :440.0	Max. :1.0390

pH	sulphates	alcohol	quality
Min. :2.720	Min. :0.2200	Min. : 8.00	Min. :3.000
1st Qu.:3.090	1st Qu.:0.4100	1st Qu.: 9.50	1st Qu.:5.000
Median :3.180	Median :0.4700	Median :10.40	Median :6.000
Mean :3.188	Mean :0.4898	Mean :10.51	Mean :5.878
3rd Qu.:3.280	3rd Qu.:0.5500	3rd Qu.:11.40	3rd Qu.:6.000
Max. :3.820	Max. :1.0800	Max. :14.20	Max. :9.000

Exploratory Data Analysis I started the EDA by first examine if there is any skew in the dataset as well as seeing the general distribution of the 2. It is interesting to further look into the dataset since there are so many variables and some of them will significantly impact the future analysis if we do not take them into factors.

```
set.seed(1)
library(ggplot2)
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
install.packages("corrplot")
```

The downloaded binary packages are in

/var/folders/kj/zb0ctvqs5zjbr\_x5cj8crxjh0000gn/T//RtmpmoMCiW/downloaded\_packages

```
library(corrplot)
```

corrplot 0.92 loaded

```
visualize_distribution <- function(df, dataset_name) {
  input_vars <- names(df)[1:11]
  hist_plots <- lapply(input_vars, function(var) {
    ggplot(df, aes_string(x = var)) +
```

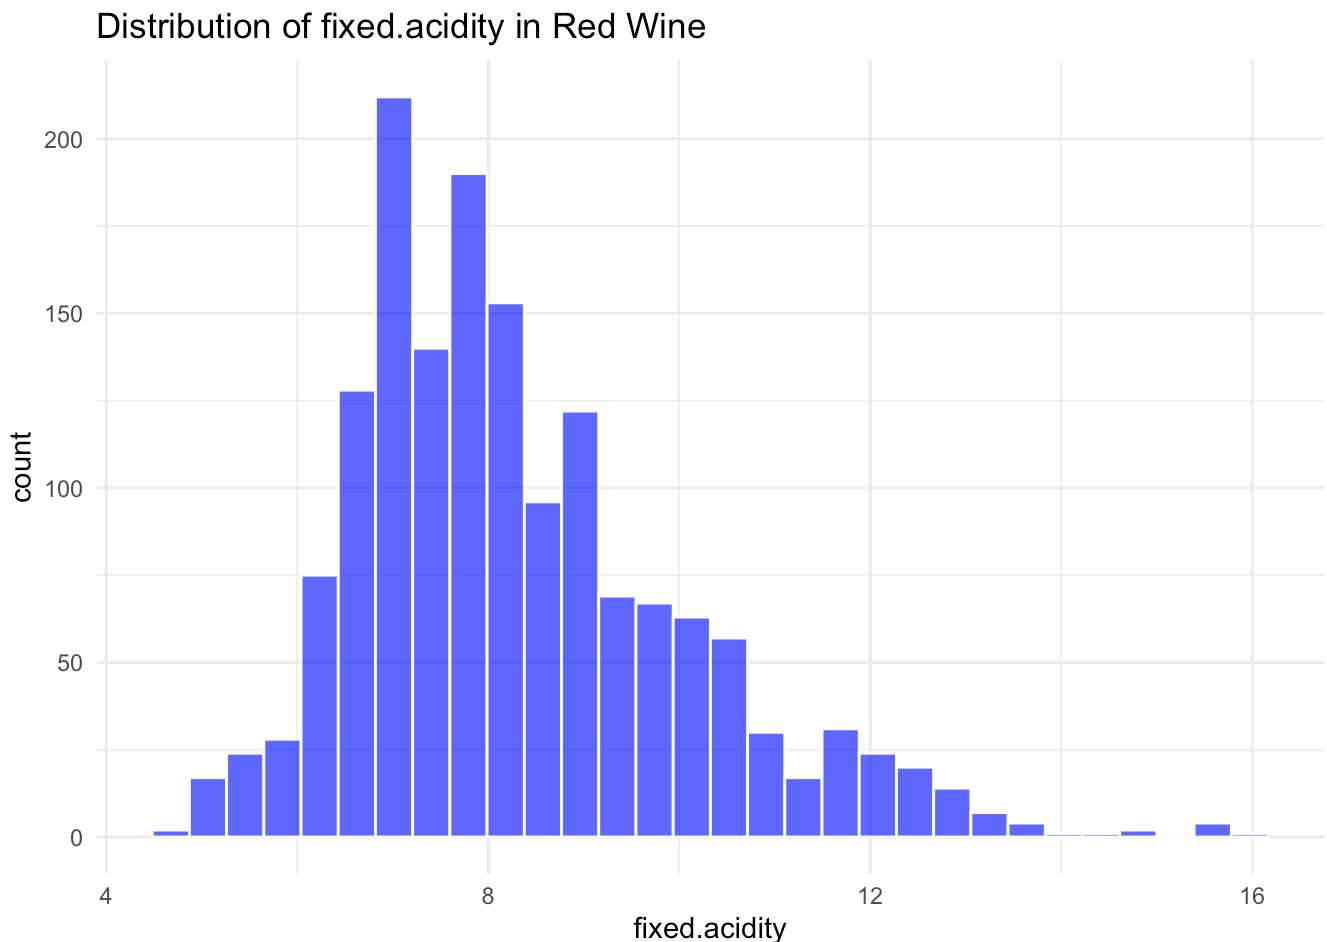
```
geom_histogram(bins = 30, fill = "blue", color = "white", alpha = 0.7) +
labs(title = paste("Distribution of", var, "in", dataset_name), x = var) +
theme_minimal()
})
box_plots <- lapply(input_vars, function(var) {
  ggplot(df, aes_string(x = "factor(quality)", y = var)) +
  geom_boxplot() +
  labs(title = paste("Boxplot of", var, "by quality in", dataset_name), x = "Quality")
  theme_minimal()
})
return(list(histograms = hist_plots, boxplots = box_plots))
}
red_wine_visuals <- visualize_distribution(red_wine, "Red Wine")
```

Warning: `aes\_string()` was deprecated in ggplot2 3.0.0.

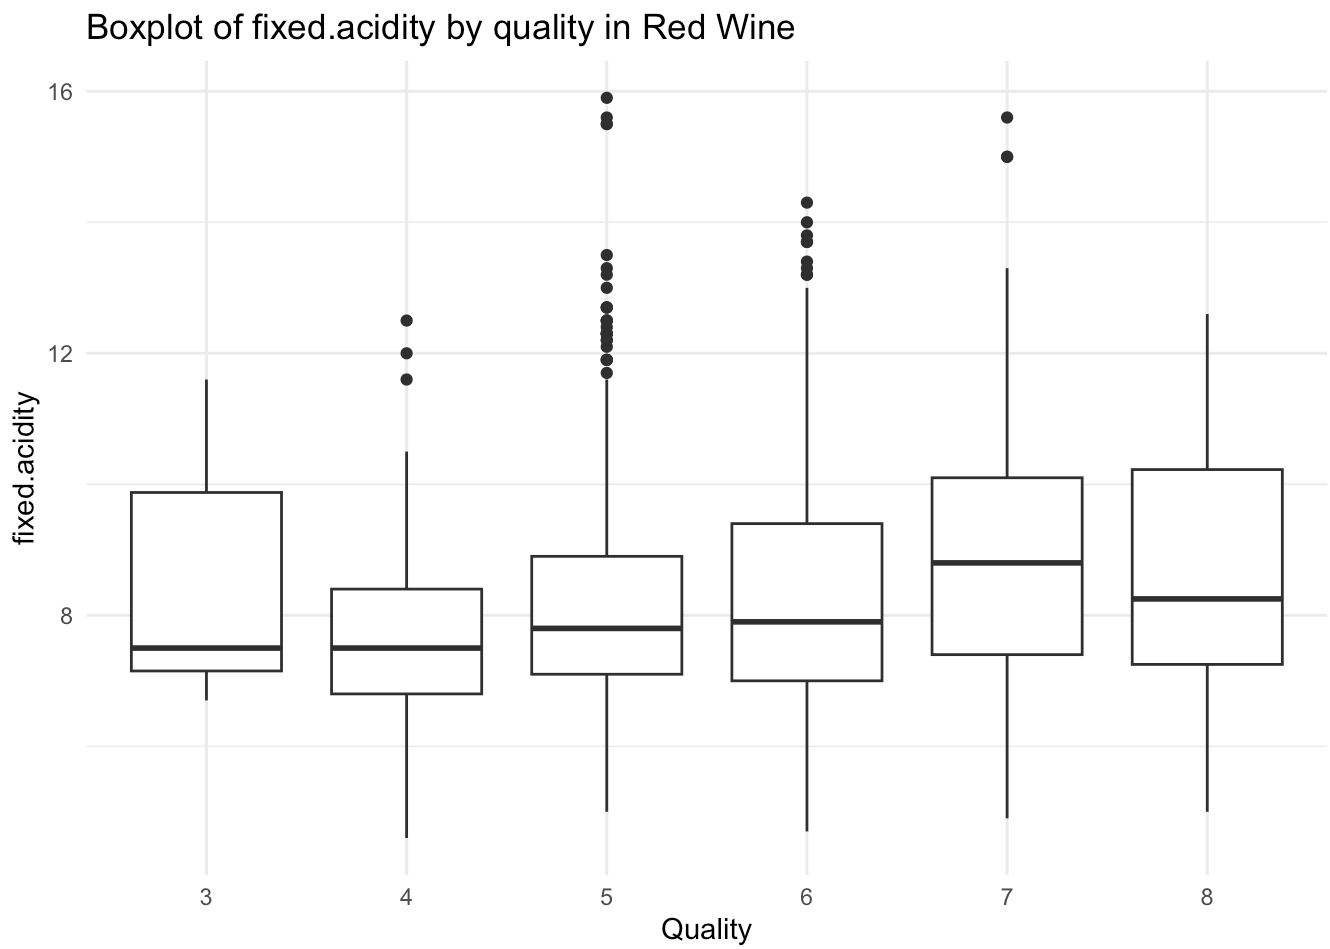
• Please use tidy evaluation idioms with `aes()`.

• See also `vignette("ggplot2-in-packages")` for more information.

```
white_wine_visuals <- visualize_distribution(white_wine, "White Wine")
print(red_wine_visuals$histograms[[1]])
```

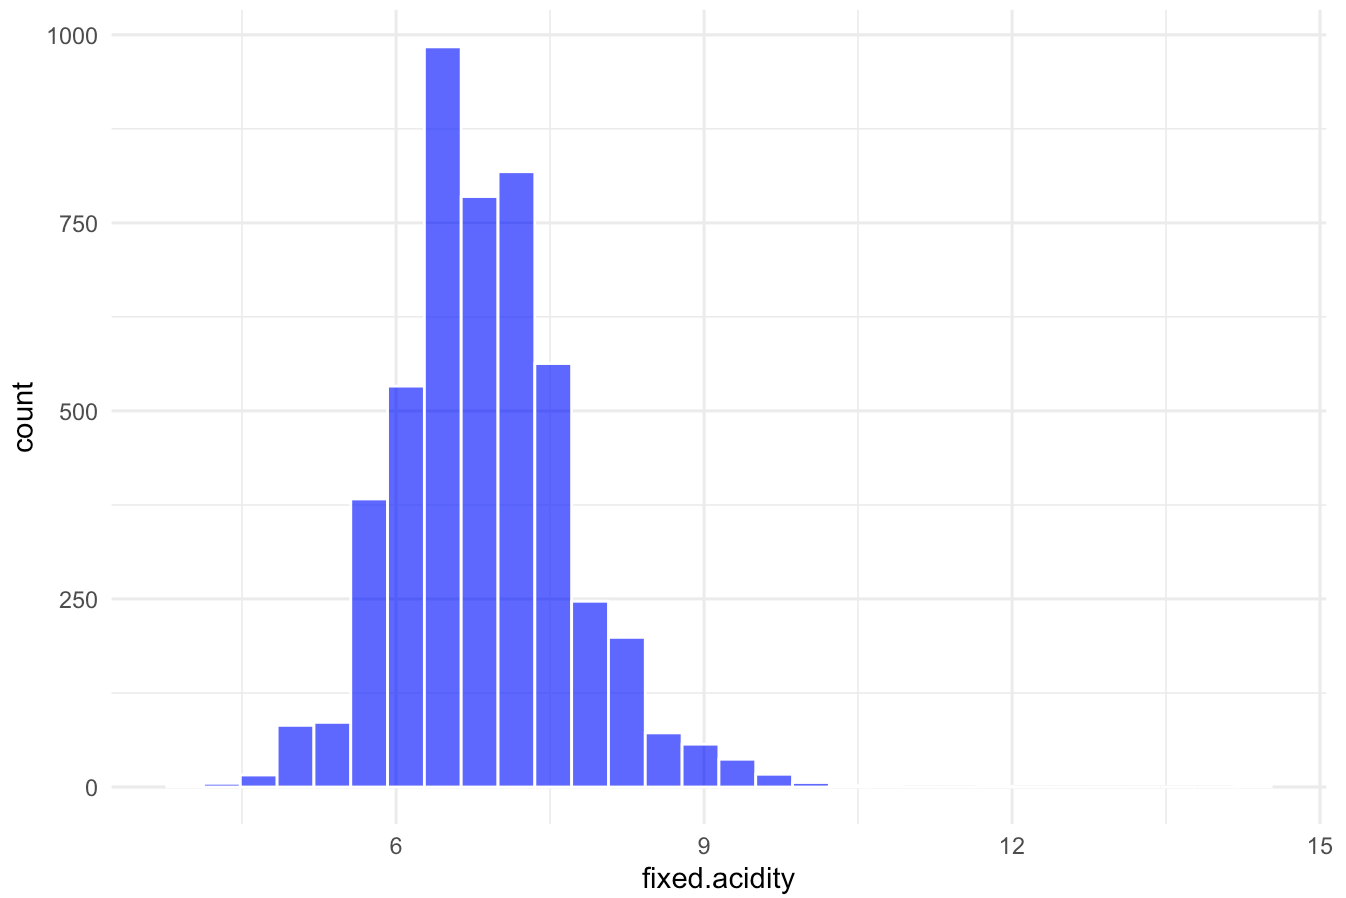


```
print(red_wine_visuals$boxplots[[1]])
```




```
print(white_wine_visuals$histograms[[1]])
```

## Distribution of fixed.acidity in White Wine




```
print(white_wine_visuals$boxplots[[1]])
```

Boxplot of fixed.acidity by quality in White Wine




```
corr_red <- cor(red_wine[, 1:11])  
corr_white <- cor(white_wine[, 1:11])  
corrplot(corr_red, method = "circle", title = "Correlation Matrix - Red Wine", mar = c(0,
```

Correlation Matrix - Red Wine



```
corrplot(corr_white, method = "circle", title = "Correlation Matrix - White Wine", mar =
```

## Correlation Matrix - White Wine



```
table(red_wine$quality)
```

```
3  4  5  6  7  8
10 53 681 638 199 18
```

```
table(white_wine$quality)
```

```
3  4  5  6  7  8  9
20 163 1457 2198 880 175 5
```

This part is mainly about examine the correlation among the variables and how they will impact the outcome of our predictions. We also do some initial analysis by plotting the graph by counting the frequency for the variables, just to have some grasps about how the data perform in general or if there is any outliers for both dataset.

```
library(ggplot2)
library(dplyr)
library(caret)
```

Loading required package: lattice



```
install.packages("ggcorrplot")
```

The downloaded binary packages are in

/var/folders/kj/zb0ctvqs5zjbr\_x5cj8crxjh0000gn/T//RtmpmoMCiW/downloaded\_packages

```
library(ggcorrplot)
library(GGally)
```

Registered S3 method overwritten by 'GGally':

```
method from
+.gg      ggplot2
```

```
library(cluster)
install.packages("factoextra")
```

The downloaded binary packages are in

/var/folders/kj/zb0ctvqs5zjbr\_x5cj8crxjh0000gn/T//RtmpmoMCiW/downloaded\_packages

```
library(factoextra)
```

Welcome! Want to learn more? See two factoextra-related books at <https://goo.gl/ve3WBa>

```
library(cowplot)
library(randomForest)
```

randomForest 4.7-1.1

Type rfNews() to see new features/changes/bug fixes.

Attaching package: 'randomForest'

The following object is masked from 'package:dplyr':

combine

The following object is masked from 'package:ggplot2':

margin


```
detect_outliers <- function(df, var) {
  Q1 <- quantile(df[[var]], 0.25)
  Q3 <- quantile(df[[var]], 0.75)
  IQR <- Q3 - Q1
  outliers <- df |>
    filter(df[[var]] < (Q1 - 1.5 * IQR) | df[[var]] > (Q3 + 1.5 * IQR))
}
```

```

    return(outliers)
}
red_outliers <- lapply(names(red_wine)[1:11], detect_outliers, df = red_wine)
white_outliers <- lapply(names(white_wine)[1:11], detect_outliers, df = white_wine)
set.seed(123)
red_rf <- randomForest(quality ~ ., data = red_wine, importance = TRUE)
white_rf <- randomForest(quality ~ ., data = white_wine, importance = TRUE)
red_imp <- varImpPlot(red_rf, main = "Red Wine - Feature Importance")

```

## Red Wine - Feature Importance




```

white_imp <- varImpPlot(white_rf, main = "White Wine - Feature Importance")

```

## White Wine - Feature Importance



```
pca_visualization <- function(df, dataset_name) {
  pca <- prcomp(df[, 1:11], scale. = TRUE)
  pca_df <- data.frame(pca$x)
  pca_df$quality <- df$quality

  ggplot(pca_df, aes(PC1, PC2, color = factor(quality))) +
    geom_point(alpha = 0.7) +
    labs(title = paste("PCA Clustering for", dataset_name), color = "Quality") +
    theme_minimal()
}

red_pca_plot <- pca_visualization(red_wine, "Red Wine")
white_pca_plot <- pca_visualization(white_wine, "White Wine")
```

In the graphs above, I mainly want to explore how different variables can affect the performance if we do model analysis as it is a very significant factor to consider. And the graph clearly demonstrates the the importance of each variables.

```
set.seed(1)
# Perform best subset selection for white wines
regfit.full_white <- regsubsets(quality ~ ., data = white_wine, nvmax = 12)
reg.summary_white <- summary(regfit.full_white)
(reg.summary_white)
```

Subset selection object

Call: regsubsets.formula(quality ~ ., data = white\_wine, nvmax = 12)

11 Variables (and intercept)

	Forced in	Forced out
fixed.acidity	FALSE	FALSE
volatile.acidity	FALSE	FALSE
citric.acid	FALSE	FALSE
residual.sugar	FALSE	FALSE
chlorides	FALSE	FALSE
free.sulfur.dioxide	FALSE	FALSE
total.sulfur.dioxide	FALSE	FALSE
density	FALSE	FALSE
pH	FALSE	FALSE
sulphates	FALSE	FALSE
alcohol	FALSE	FALSE

1 subsets of each size up to 11

Selection Algorithm: exhaustive

	fixed.acidity	volatile.acidity	citric.acid	residual.sugar	chlorides
1 ( 1 )	" "	" "	" "	" "	" "
2 ( 1 )	" "	"*"	" "	" "	" "
3 ( 1 )	" "	"*"	" "	"*"	" "
4 ( 1 )	" "	"*"	" "	"*"	" "
5 ( 1 )	" "	"*"	" "	"*"	" "
6 ( 1 )	" "	"*"	" "	"*"	" "
7 ( 1 )	" "	"*"	" "	"*"	" "
8 ( 1 )	"*"	"*"	" "	"*"	" "
9 ( 1 )	"*"	"*"	" "	"*"	" "
10 ( 1 )	"*"	"*"	" "	"*"	"*"
11 ( 1 )	"*"	"*"	"*"	"*"	"*"

	free.sulfur.dioxide	total.sulfur.dioxide	density	pH	sulphates
1 ( 1 )	" "	" "	" "	" "	" "
2 ( 1 )	" "	" "	" "	" "	" "
3 ( 1 )	" "	" "	" "	" "	" "
4 ( 1 )	"*"	" "	" "	" "	" "
5 ( 1 )	" "	" "	"*"	"*"	" "
6 ( 1 )	" "	" "	"*"	"*"	"*"
7 ( 1 )	"*"	" "	"*"	"*"	"*"
8 ( 1 )	"*"	" "	"*"	"*"	"*"
9 ( 1 )	"*"	"*"	"*"	"*"	"*"
10 ( 1 )	"*"	"*"	"*"	"*"	"*"
11 ( 1 )	"*"	"*"	"*"	"*"	"*"

	alcohol
1 ( 1 )	"*"
2 ( 1 )	"*"
3 ( 1 )	"*"
4 ( 1 )	"*"
5 ( 1 )	"*"
6 ( 1 )	"*"
7 ( 1 )	"*"
8 ( 1 )	"*"

```
9 ( 1 ) "*"
10 ( 1 ) "*"
11 ( 1 ) "*"
```

```
(reg.summary_white$adjr2)
```

```
[1] 0.1895598 0.2399208 0.2580716 0.2633925 0.2703282 0.2757705 0.2790891
[8] 0.2805767 0.2805130 0.2803931 0.2802536
```

```
# Find best subset based on BIC
best_subset_white <- which.min(reg.summary_white$bic)
# Perform best subset selection for red wines
regfit.full_red <- regsubsets(quality ~ ., data = red_wine, nvmax = 12)
reg.summary_red <- summary(regfit.full_red)
(reg.summary_red)
```

Subset selection object

Call: regsubsets.formula(quality ~ ., data = red\_wine, nvmax = 12)

11 Variables (and intercept)

	Forced in	Forced out
fixed.acidity	FALSE	FALSE
volatile.acidity	FALSE	FALSE
citric.acid	FALSE	FALSE
residual.sugar	FALSE	FALSE
chlorides	FALSE	FALSE
free.sulfur.dioxide	FALSE	FALSE
total.sulfur.dioxide	FALSE	FALSE
density	FALSE	FALSE
pH	FALSE	FALSE
sulphates	FALSE	FALSE
alcohol	FALSE	FALSE

1 subsets of each size up to 11

Selection Algorithm: exhaustive

	fixed.acidity	volatile.acidity	citric.acid	residual.sugar	chlorides
1 ( 1 )	" "	" "	" "	" "	" "
2 ( 1 )	" "	"*"	" "	" "	" "
3 ( 1 )	" "	"*"	" "	" "	" "
4 ( 1 )	" "	"*"	" "	" "	" "
5 ( 1 )	" "	"*"	" "	" "	"*"
6 ( 1 )	" "	"*"	" "	" "	"*"
7 ( 1 )	" "	"*"	" "	" "	"*"
8 ( 1 )	" "	"*"	"*"	" "	"*"
9 ( 1 )	" "	"*"	"*"	"*"	"*"
10 ( 1 )	"*"	"*"	"*"	"*"	"*"
11 ( 1 )	"*"	"*"	"*"	"*"	"*"

	free.sulfur.dioxide	total.sulfur.dioxide	density	pH	sulphates
1 ( 1 )	" "	" "	" "	" "	" "
2 ( 1 )	" "	" "	" "	" "	" "
3 ( 1 )	" "	" "	" "	" "	"*"

4	( 1 )	" "	"*"	" "	" " "	"*"
5	( 1 )	" "	"*"	" "	" " "	"*"
6	( 1 )	" "	"*"	" "	"*"	"*"
7	( 1 )	"*"	"*"	" "	"*"	"*"
8	( 1 )	"*"	"*"	" "	"*"	"*"
9	( 1 )	"*"	"*"	" "	"*"	"*"
10	( 1 )	"*"	"*"	" "	"*"	"*"
11	( 1 )	"*"	"*"	"*"	"*"	"*"

alcohol

1	( 1 )	"*"
2	( 1 )	"*"
3	( 1 )	"*"
4	( 1 )	"*"
5	( 1 )	"*"
6	( 1 )	"*"
7	( 1 )	"*"
8	( 1 )	"*"
9	( 1 )	"*"
10	( 1 )	"*"
11	( 1 )	"*"

```
(reg.summary_red$adjr2)
```


```
[1] 0.2262502 0.3161465 0.3346482 0.3421357 0.3494588 0.3547509 0.3566527
[8] 0.3567060 0.3565489 0.3562479 0.3561195
```

```
# Find best subset based on BIC
best_subset_red <- which.min(reg.summary_red$bic)
```


Here we are just running the best subset selection so that we can run the future model using this as a reference. It is also important to note that we are looking for the one that has the smallest BIC is the set of predictors that we are going to use.

```
set.seed(1)
par(mfrow = c(1,2))
plot(regfit.full_red, scale = "bic", main = 'Best subset for red wines')
plot(regfit.full_white, scale = "bic", main = 'Best subset for white wines')
```

## Best subset for red wines



## Best subset for white wines



In this case, we are just graphing the best subset based on the result above.

```
library(caret)
library(randomForest)
set.seed(1)
adjusted_r2 <- function(r2, n, p) {
  return(1 - ((1 - r2) * (n - 1) / (n - p - 1)))
}
aic_calculator <- function(n, mse, num_params) {
  return(n * log(mse) + 2 * num_params)
}
ctrl <- trainControl(method = "cv", number = 10)
model_white_rf <- train(
  quality ~ volatile.acidity + citric.acid + chlorides + free.sulfur.dioxide +
    total.sulfur.dioxide + pH + sulphates + alcohol,
  data = white_wine,
  method = "rf",
  trControl = ctrl
)
white_rf_results <- model_white_rf$results
best_white_rf <- model_white_rf$bestTune
final_model_white_rf <- model_white_rf$finalModel
white_r2 <- max(white_rf_results$Rsquared)
n_white <- nrow(white_wine)
p_white <- ncol(white_wine) - 1
```

```
white_adj_r2 <- adjusted_r2(white_r2, n_white, p_white)
white_rf_predictions <- predict(final_model_white_rf, newdata = white_wine)
white_mse <- mean((white_wine$quality - white_rf_predictions)^2)
white_aic <- aic_calculator(n_white, white_mse, p_white)
model_red_rf <- train(
  quality ~ volatile.acidity + citric.acid + chlorides + free.sulfur.dioxide +
    total.sulfur.dioxide + pH + sulphates + alcohol,
  data = red_wine,
  method = "rf",
  trControl = ctrl
)
red_rf_results <- model_red_rf$results
best_red_rf <- model_red_rf$bestTune
final_model_red_rf <- model_red_rf$finalModel
red_r2 <- max(red_rf_results$Rsquared)
n_red <- nrow(red_wine)
p_red <- ncol(red_wine) - 1
red_adj_r2 <- adjusted_r2(red_r2, n_red, p_red)
red_rf_predictions <- predict(final_model_red_rf, newdata = red_wine)
red_mse <- mean((red_wine$quality - red_rf_predictions)^2)
red_aic <- aic_calculator(n_red, red_mse, p_red)
cat("Best Random Forest model for white wine quality:\n")
```

Best Random Forest model for white wine quality:

```
print(best_white_rf)
```

```
mtry
2    5
```

```
cat("\nWhite wine performance metrics:\n")
```

White wine performance metrics:

```
cat("Adjusted R^2:", white_adj_r2, "\n")
```

Adjusted R<sup>2</sup>: 0.5408396

```
cat("MSE:", white_mse, "\n")
```

MSE: 0.0671829

```
cat("AIC:", white_aic, "\n")
```

AIC: -13204.25



```
cat("\nBest Random Forest model for red wine quality:\n")
```

Best Random Forest model for red wine quality:

```
print(best_red_rf)
```

```
mtry  
1    2
```

```
cat("\nRed wine performance metrics:\n")
```

Red wine performance metrics:

```
cat("Adjusted R^2:", red_adj_r2, "\n")
```

Adjusted R<sup>2</sup>: 0.5156118

```
cat("MSE:", red_mse, "\n")
```

MSE: 0.0698625

```
cat("AIC:", red_aic, "\n")
```

AIC: -4233.301

Lastly, I performed the random forest model to train the model so that we can get the best adjusted r square, mse, aic. We can tell that from the result, it is probably not the best model to run for this specific dataset since a lot of the performance metrics are extremely poor.